

Submission Date: **5.11.2025**

1 Bias-Variance Trade-off

You are working as a data scientist for a very important company and, suddenly, you are tasked to train a *polynomial* model on the data from Figure 1.




Figure 1: Noisy Data.

1. Only from looking at the plotted data, which degree of a polynomial would be your best guess that would fit the data without over- or underfitting it? Why?
2. Predict how the bias and variance (high / low) and fit (under / over) would be for:
 - A constant model (i.e. a horizontal line)
 - A linear model
 - A polynomial model of degree 10

 **Using Python and SciKit-Learn:** Included in the exercise are two files (`data_train.csv` and `data_test.csv`) containing the noisy data in a csv format, where the first column represents the input features x and the second column contains the measured outputs y . Additionally, there is a jupyter notebook that contains some code snippets to help you get started. The notebook requires a virtual environment using anaconda or venv with numpy, sklearn, jupyter and matplotlib.

3. Use `numpy` to load the two splits of the dataset and use `matplotlib` to try to recreate the plot from Figure 1.
4. Using the `sklearn.linear_model.LinearRegression` class, fit a linear model to the data in the `data_train.csv` file. Compute the MSE error

for the data in `data_test.csv` and plot the predicted line together with the original data.

5. Create a class that can fit a polynomial model of an arbitrary degree. For example `m = PolynomialModel(degree=23)`. Include a `fit(x_train, y_train)` method for fitting some training data and a `predict(x)` method for performing inference for the values of unseen data. Using your class, fit models with $degrees \in \{0, 1, 2, 3, 4, 5, 20\}$ and plot the fitted curves along with the original data.

Hint: You can use `sklearn.linear_model.LinearRegression` and `sklearn.preprocessing.PolynomialFeatures` together to define non-linear models by combining a linear model with basis functions.

6. Create a function `measure_bias_var_mse(model, x_train, y_train, x_test, y_test, n)` that estimates the bias, variance and, mean square error of different models.

Hints: The function should:

- (a) sample *with replacement* n datasets $\{\mathcal{D}_1^{(train)}, \dots, \mathcal{D}_n^{(train)}\}$ from the training dataset,
- (b) fit a model for each re-sampled dataset $f_i(x; \mathcal{D}_i^{(train)})$
- (c) generate predictions over the test data with each model and store the predictions in some array $pred_{i,j} = \hat{f}(x_i^{(test)}; \mathcal{D}_j^{(train)})$
- (d) compute the *expected prediction model*:

$$\begin{aligned}\bar{f}(x^{(test)}) &= \underset{\mathcal{D}^{(test)} \sim \mathcal{P}^n}{\mathbb{E}} [\hat{f}(x^{(test)}; \mathcal{D})] \\ &= \frac{1}{n} \sum_{i=1}^n \hat{f}(x^{(test)}; \mathcal{D}_i^{(train)})\end{aligned}\quad (1)$$

- (e) compute the *bias* using the expected prediction model:

$$\begin{aligned}\text{bias} &= \underset{x,y}{\mathbb{E}} \left[(\bar{f}(x^{(test)}) - y^{(test)})^2 \right] \\ &= \frac{1}{|\mathcal{D}^{(test)}|} \sum_{i=1}^{|\mathcal{D}^{(test)}|} (\bar{f}(x_i^{(test)}) - y_i^{(test)})^2\end{aligned}\quad (2)$$

- (f) compute the *variance* using the expected prediction model:

$$\begin{aligned}\text{variance} &= \underset{x,\mathcal{D}}{\mathbb{E}} \left[(\hat{f}(x; \mathcal{D}) - \bar{f}(x)) \right]^2 \\ &= \frac{1}{n \cdot |\mathcal{D}^{(test)}|} \sum_{i=1}^{|\mathcal{D}^{(test)}|} \sum_{j=1}^n (\bar{f}(x_i^{(test)}) - \hat{f}(x_i^{(test)}; \mathcal{D}_j^{(train)}))^2\end{aligned}\quad (3)$$

(g) compute the *average mean square error*:

$$\text{MSE} = \frac{1}{n \cdot |\mathcal{D}^{\text{test}}|} \sum_{i=1}^{|\mathcal{D}^{\text{test}}|} \sum_{j=1}^n (\hat{f}(x_i^{\text{test}}; \mathcal{D}_j^{\text{train}}) - y_i^{\text{(test)}})^2 \quad (4)$$

7. Using your function for estimating the MSE, bias and variance, compute them for models with $\text{degrees} \in \{0, 1, \dots, 10\}$. Generate a plot of the MSE, bias and variance, where the x axis represents the degrees of the polynomial model, and the y axis are the estimated metrics. Do the computed values and plot match your predictions from question 2.?